

model

December 11, 2025

```
[8]: # =====
# Final Project: Train neural net on MNIST and evaluate on local digits
# =====

DEMO = True

import os
import glob
from io import BytesIO

import numpy as np
import matplotlib.pyplot as plt
from PIL import Image, ImageFilter, ImageEnhance

import torch
from torch import nn
import torch.nn.functional as F
from torchvision import datasets, transforms
from torch.utils.data import Dataset, Subset
from sklearn.metrics import classification_report, confusion_matrix

# -----
# Device
# -----
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
```

Using device: cpu

```
[9]: # =====
# 1. Data and augmentation setup
# =====

class RaiseDarkPoint(object):
    def __init__(self, grayRange=(10, 60)):
        self.grayRange = grayRange
```

```

def __call__(self, img):
    arr = np.array(img).astype(np.float32)
    grayVal = np.random.uniform(*self.grayRange)
    mask = arr < 30
    arr[mask] = arr[mask] + grayVal
    arr = np.clip(arr, 0, 255)
    return Image.fromarray(arr.astype(np.uint8), mode="L")

class LowerWhitePoint(object):
    def __init__(self, factorRange=(0.7, 0.95)):
        self.factorRange = factorRange

    def __call__(self, img):
        arr = np.array(img, dtype=np.float32)
        factor = float(np.random.uniform(self.factorRange[0], self.
factorRange[1]))
        arr = arr * factor
        arr = np.clip(arr, 0, 255)
        return Image.fromarray(arr.astype(np.uint8), mode="L")

class AddNoise(object):
    def __init__(self, noiseStd=0.05):
        self.noiseStd = noiseStd

    def __call__(self, img):
        arr = np.array(img).astype(np.float32)
        noise = np.random.normal(0, self.noiseStd * 255, arr.shape)
        arr = arr + noise
        arr = np.clip(arr, 0, 255)
        return Image.fromarray(arr.astype(np.uint8), mode="L")

class BubblyDigits(object):
    def __init__(self, blurRange=(0.4, 1.0), contrastRange=(1.1, 1.6)):
        self.blurRange = blurRange
        self.contrastRange = contrastRange

    def __call__(self, img):
        sigma = float(np.random.uniform(self.blurRange[0], self.blurRange[1]))
        img = img.filter(ImageFilter.GaussianBlur(radius=sigma))
        c = float(np.random.uniform(self.contrastRange[0], self.
contrastRange[1]))
        img = ImageEnhance.Contrast(img).enhance(c)
        return img

```

```

class JPEGCompression(object):
    def __init__(self, qualityRange=(40, 80)):
        self.qualityRange = qualityRange

    def __call__(self, img):
        q = int(np.random.randint(self.qualityRange[0], self.qualityRange[1]))
        buf = BytesIO()
        img.save(buf, format="JPEG", quality=q)
        buf.seek(0)
        return Image.open(buf).convert("L")

trainTransform = transforms.Compose(
    [
        transforms.RandomApply(
            [
                transforms.Pad(4, fill=0),
                transforms.RandomCrop(28),
            ],
            p=0.5,
        ),
        transforms.RandomApply(
            [
                transforms.RandomAffine(
                    degrees=(-5, 20),
                    translate=(0.25, 0.25),
                    scale=(0.6, 1.1),
                    fill=0,
                )
            ],
            p=0.5,
        ),
        transforms.RandomApply([AddNoise()], p=0.1),
        transforms.RandomApply([BubblyDigits()], p=0.5),
        transforms.RandomApply([LowerWhitePoint()], p=0.4),
        transforms.RandomApply([RaiseDarkPoint()], p=0.4),
        transforms.RandomApply(
            [transforms.GaussianBlur(kernel_size=3, sigma=(0.1, 1.5))], p=1
        ),
        transforms.RandomApply([JPEGCompression()], p=0.25),
        transforms.ToTensor(),
        transforms.Normalize((0.5,), (0.5,)),
    ]
)

transform = transforms.Compose(

```

```

        [transforms.ToTensor(), transforms.Normalize((0.5,), (0.5,))]
    )

def PreprocessSingle(image):
    if image.mode != "RGB":
        image = image.convert("RGB")

    blurred = image.filter(ImageFilter.GaussianBlur(radius=0.75))
    arr = np.array(blurred).astype(np.float32)
    gray = arr.mean(axis=2)

    minVal = gray.min()
    maxVal = gray.max()

    if maxVal - minVal < 1e-6:
        stretched = np.zeros_like(gray)
    else:
        stretched = (gray - minVal) / (maxVal - minVal) * 255.0

    stretched = stretched.astype(np.uint8)
    return Image.fromarray(stretched, mode="L")

# -----
# MNIST train/val/test split
# -----
fullTrainAug = datasets.MNIST(
    "~/pytorch/MNIST_data/", download=True, train=True,
    transform=trainTransform
)
fullTrainPlain = datasets.MNIST(
    "~/pytorch/MNIST_data/", download=True, train=True, transform=transform
)

trainSize = int(0.9 * len(fullTrainAug))
valSize = len(fullTrainAug) - trainSize

indices = torch.randperm(len(fullTrainAug)).tolist()
trainIdx = indices[:trainSize]
valIdx = indices[trainSize:]

trainSet = Subset(fullTrainAug, trainIdx)      # with augmentation
valSet = Subset(fullTrainPlain, valIdx)         # no augmentation

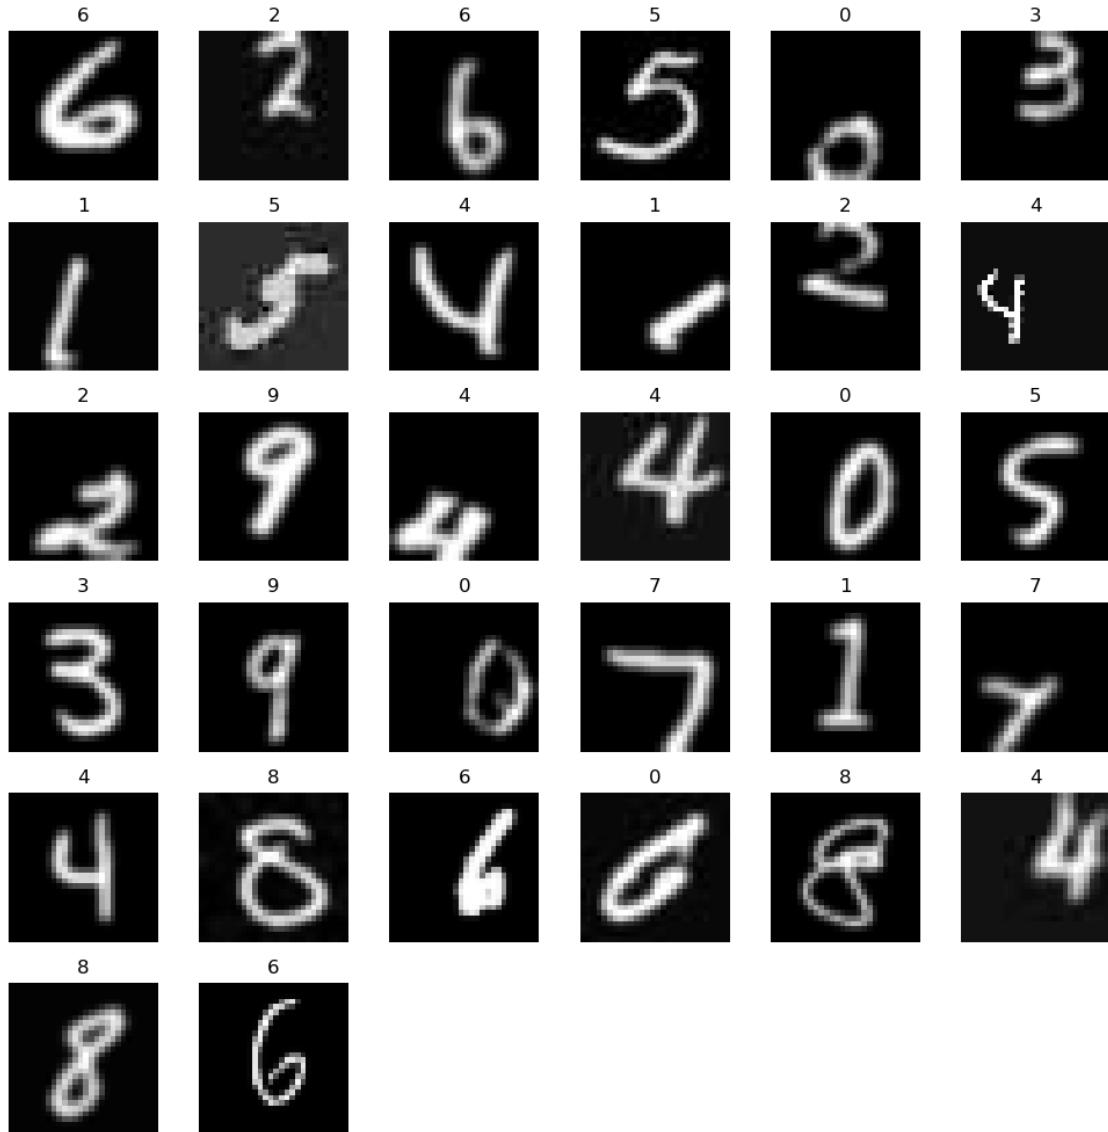
trainLoader = torch.utils.data.DataLoader(trainSet, batch_size=64, shuffle=True)
valLoader = torch.utils.data.DataLoader(valSet, batch_size=64, shuffle=False)

```

```
testSet = datasets.MNIST(  
    "~/pytorch/MNIST_data/", download=True, train=False, transform=transform  
)  
testLoader = torch.utils.data.DataLoader(testSet, batch_size=64, shuffle=False)  
  
print("Trainloader loaded:", len(trainLoader))  
print("Valloader loaded:", len(valLoader))  
  
dataIter = iter(trainLoader)  
imgs, labels = next(dataIter)  
imgsVis = imgs * 0.5 + 0.5  
  
plt.figure(figsize=(10, 10))  
for i in range(32):  
    plt.subplot(6, 6, i + 1)  
    plt.imshow(imgsVis[i].squeeze().cpu(), cmap="gray")  
    plt.title(labels[i].item())  
    plt.axis("off")  
plt.tight_layout()  
plt.show(block=False)  
plt.pause(0.001)
```

Trainloader loaded: 844

Valloader loaded: 94



```
[10]: # =====
# 2. Model definition
# =====
```

```
class MNISTMLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 512)
        self.bn1 = nn.BatchNorm1d(512)
        self.fc2 = nn.Linear(512, 256)
        self.bn2 = nn.BatchNorm1d(256)
```

```

        self.fc3 = nn.Linear(256, 128)
        self.bn3 = nn.BatchNorm1d(128)
        self.fc4 = nn.Linear(128, 10)
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        x = x.view(x.shape[0], -1)
        x = F.relu(self.bn1(self.fc1(x)))
        x = self.dropout(x)
        x = F.relu(self.bn2(self.fc2(x)))
        x = self.dropout(x)
        x = F.relu(self.bn3(self.fc3(x)))
        x = self.fc4(x)
        return x

model = MNISTMLP().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=5e-4, weight_decay=1e-4)

scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    factor=0.5,
    patience=5,
)

print(model)

```

```

MNISTMLP(
  (fc1): Linear(in_features=784, out_features=512, bias=True)
  (bn1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (fc2): Linear(in_features=512, out_features=256, bias=True)
  (bn2): BatchNorm1d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (fc3): Linear(in_features=256, out_features=128, bias=True)
  (bn3): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (fc4): Linear(in_features=128, out_features=10, bias=True)
  (dropout): Dropout(p=0.3, inplace=False)
)

```

```
[11]: # =====
# 3. Training + validation loop
# =====

if not DEMO:
```

```

numEpochs = 100
epochTrainLosses = []
valLosses = []
learningRates = []

for epoch in range(numEpochs):
    model.train()
    runningLoss = 0.0

    for images, labels in trainLoader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        logits = model(images)
        loss = criterion(logits, labels)
        loss.backward()
        optimizer.step()

        runningLoss += loss.item()

    epochLoss = runningLoss / len(trainLoader)
    epochTrainLosses.append(epochLoss)

    # ---- validation loss (no gradient, no augmentation) ----
    model.eval()
    valRunningLoss = 0.0
    with torch.no_grad():
        for images, labels in valLoader:
            images, labels = images.to(device), labels.to(device)
            logits = model(images)
            loss = criterion(logits, labels)
            valRunningLoss += loss.item()

    valLoss = valRunningLoss / len(valLoader)
    valLosses.append(valLoss)

    # Step scheduler on validation loss, then read scalar LR
    scheduler.step(valLoss)
    currentLr = optimizer.param_groups[0]["lr"]
    learningRates.append(currentLr)

    print(
        f"Epoch {epoch+1}/{numEpochs}, "
        f"Train Loss: {epochLoss:.4f}, Val Loss: {valLoss:.4f} "
        f"Learning Rate: {currentLr:.6f}"
    )

```

```

print("Training complete.")

epochs = range(1, numEpochs + 1)

plt.figure(figsize=(10, 4))
plt.plot(epochs, epochTrainLosses, label="Train Loss")
plt.plot(epochs, valLosses, label="Val Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Train and Validation Loss per Epoch")
plt.legend()
plt.tight_layout()
plt.show(block=False)
plt.pause(0.001)

# Learning rate per epoch
plt.figure(figsize=(10, 4))
plt.plot(epochs, learningRates, label="Learning Rate")
plt.xlabel("Epoch")
plt.ylabel("LR")
plt.title("Learning Rate per Epoch")
plt.legend()
plt.tight_layout()
plt.show(block=False)
plt.pause(0.001)

torch.save(model.state_dict(), "model.pth")
else:
    # DEMO mode: load pretrained model
    model.load_state_dict(torch.load("model.pth", map_location=device))
    model.eval()
    print("Loaded model.pth (DEMO mode)")

```

Loaded model.pth (DEMO mode)

[12]:

```

# =====
# 4. Evaluation helpers
# =====

def PlotConfusionMatrix(cm, title):
    plt.figure(figsize=(7, 6))
    plt.imshow(cm, interpolation="nearest", cmap="Blues")
    plt.title(f"{title} Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("True")

```

```

plt.xticks(range(10))
plt.yticks(range(10))
plt.colorbar()
for r in range(cm.shape[0]):
    for c in range(cm.shape[1]):
        v = cm[r, c]
        if v != 0:
            plt.text(c, r, str(v), ha="center", va="center", fontsize=9)
plt.tight_layout()
plt.show(block=False)
plt.pause(0.001)

def PlotMissesByDigit(allPreds, allLabels, title):
    misses = []
    for d in range(10):
        misses.append(int(np.sum((allLabels == d) & (allPreds != d)))))

    plt.figure(figsize=(8, 4))
    plt.bar(range(10), misses)
    plt.title(f"{title} Misses by True Digit")
    plt.xlabel("True Digit")
    plt.ylabel("Miss Count")
    plt.xticks(range(10))
    plt.tight_layout()
    plt.show(block=False)
    plt.pause(0.001)

def PlotGroupAccuracy(allPreds, allLabels, allGroupIDs, title):
    valid = allGroupIDs >= 0
    if not np.any(valid):
        return

    groupAcc = {}
    for gid in np.unique(allGroupIDs[valid]):
        mask = (allGroupIDs == gid)
        groupAcc[gid] = float(np.mean(allPreds[mask] == allLabels[mask]))

    groups = list(groupAcc.keys())
    accs = [groupAcc[g] for g in groups]

    plt.figure(figsize=(10, 5))
    plt.bar([str(g) for g in groups], accs)
    plt.ylim(0, 1.05)
    plt.xlabel("Group ID")
    plt.ylabel("Accuracy")
    plt.title(f"{title} Accuracy by Group")
    plt.tight_layout()

```

```

plt.show(block=False)
plt.pause(0.001)

def PrintMispredictions(allPreds, allLabels, allPaths, title, maxMisses=None):
    missIdx = np.where(allPreds != allLabels)[0]
    print(f"\n{title} mispredictions: {len(missIdx)} / {len(allLabels)}")

    if len(allPaths) != len(allLabels):
        print("(No file paths available to print per-image misses.)")
        return

    shown = 0
    for i in missIdx:
        print(f"- {allPaths[i]} | true={int(allLabels[i])} |"
              f"pred={int(allPreds[i])}")
        shown += 1
        if maxMisses is not None and shown >= maxMisses:
            print(f"... truncated at {maxMisses} misses")
            break

def EvaluateWithDetails(loader, name="Dataset", plot=True, printMisses=True, ↴
    ↴maxMisses=None):
    model.eval()
    allPreds, allLabels = [], []
    allGroupIDs = []
    allPaths = []

    with torch.no_grad():
        for data in loader:
            if len(data) == 4:
                images, labels, group_ids, paths = data
                allGroupIDs.extend(group_ids.cpu().numpy())
                allPaths.extend(list(paths))
            elif len(data) == 3:
                images, labels, group_ids = data
                allGroupIDs.extend(group_ids.cpu().numpy())
            elif len(data) == 2:
                images, labels = data
            else:
                continue

            images, labels = images.to(device), labels.to(device)
            logits = model(images)
            _, preds = torch.max(logits, 1)

```

```

        allPreds.extend(preds.cpu().numpy())
        allLabels.extend(labels.cpu().numpy())

    allPreds = np.array(allPreds)
    allLabels = np.array(allLabels)
    allGroupIDs = np.array(allGroupIDs) if len(allGroupIDs) == len(allLabels) else None
    else None

    accuracy = (allPreds == allLabels).mean()
    errorRate = 1.0 - accuracy

    print(f"\n==== {name} RESULTS ====")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Error Rate: {errorRate:.4f}")
    print("\nClassification Report:")
    print(classification_report(allLabels, allPreds, digits=4))

    cm = confusion_matrix(allLabels, allPreds)
    print("Confusion Matrix:")
    print(cm)

    if plot:
        PlotConfusionMatrix(cm, name)
        PlotMissesByDigit(allPreds, allLabels, name)
        if allGroupIDs is not None and np.any(allGroupIDs >= 0):
            PlotGroupAccuracy(allPreds, allLabels, allGroupIDs, name)

    if printMisses:
        PrintMispredictions(allPreds, allLabels, allPaths, name, maxMisses=maxMisses)

    # preserve old behavior: MNIST call expects a float accuracy
    return accuracy

mnistAcc = EvaluateWithDetails(testLoader, "MNIST Test Set")
print("MNIST test accuracy (float):", float(mnistAcc))

```

```

===== MNIST Test Set RESULTS =====
Accuracy: 0.9910
Error Rate: 0.0090

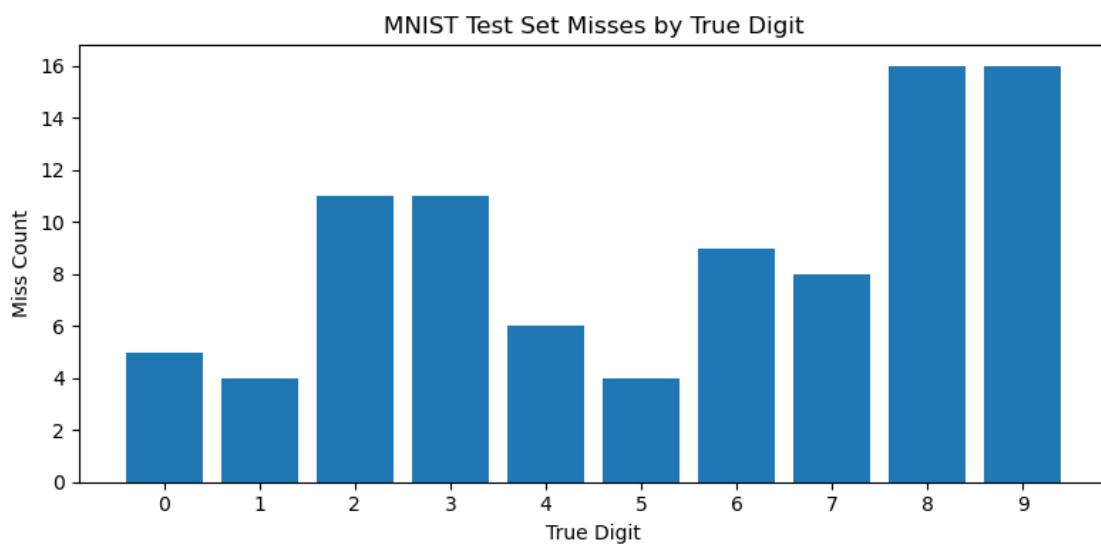
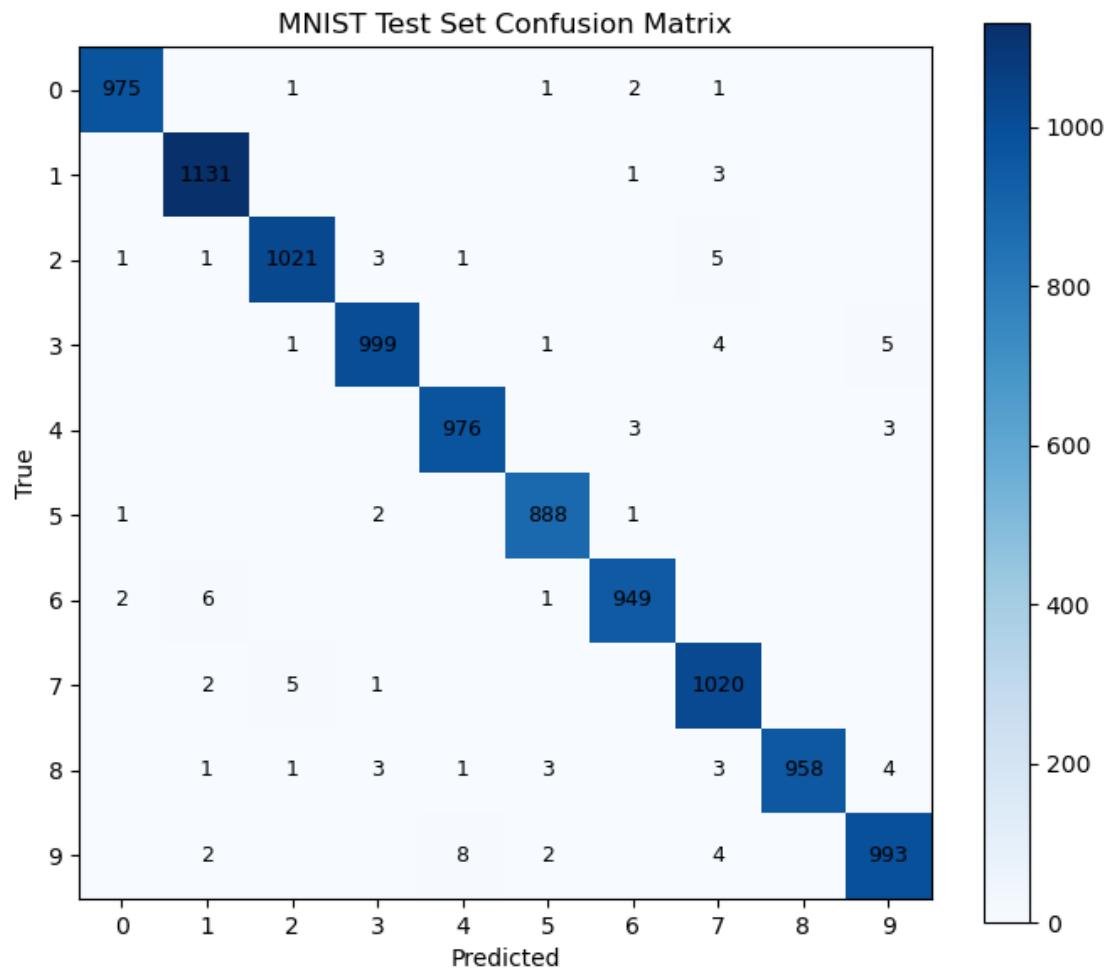
Classification Report:
      precision    recall   f1-score   support
      0       0.9959     0.9949     0.9954      980

```

1	0.9895	0.9965	0.9930	1135
2	0.9922	0.9893	0.9908	1032
3	0.9911	0.9891	0.9901	1010
4	0.9899	0.9939	0.9919	982
5	0.9911	0.9955	0.9933	892
6	0.9927	0.9906	0.9916	958
7	0.9808	0.9922	0.9865	1028
8	1.0000	0.9836	0.9917	974
9	0.9881	0.9841	0.9861	1009
accuracy			0.9910	10000
macro avg	0.9911	0.9910	0.9910	10000
weighted avg	0.9910	0.9910	0.9910	10000

Confusion Matrix:

```
[[ 975   0   1   0   0   1   2   1   0   0]
 [  0 1131   0   0   0   0   1   3   0   0]
 [  1   1 1021   3   1   0   0   5   0   0]
 [  0   0   1 999   0   1   0   4   0   5]
 [  0   0   0   0 976   0   3   0   0   3]
 [  1   0   0   2   0 888   1   0   0   0]
 [  2   6   0   0   0   1 949   0   0   0]
 [  0   2   5   1   0   0   0 1020   0   0]
 [  0   1   1   3   1   3   0   3 958   4]
 [  0   2   0   0   8   2   0   4   0 993]]
```



```
MNIST Test Set mispredictions: 90 / 10000
(No file paths available to print per-image misses.)
MNIST test accuracy (float): 0.991
```

```
[13]: # =====
# 5. Local handwritten digit dataset
# =====

class HandwrittenDigits(Dataset):
    def __init__(self, root, transform=None):
        self.paths = sorted(glob.glob(os.path.join(root, "*.png")))
        self.transform = transform
        self.labels = [int(os.path.basename(p).split("-")[0]) for p in self.
                      paths]

        # Try to parse group id like "digit-group-....png"
        self.group_ids = []
        for p in self.paths:
            parts = os.path.basename(p).split("-")
            if len(parts) > 1 and parts[1].isdigit():
                self.group_ids.append(int(parts[1]))
            else:
                self.group_ids.append(-1) # no group present

    def __len__(self):
        return len(self.paths)

    def __getitem__(self, idx):
        path = self.paths[idx]
        img = Image.open(path)
        img = PreprocessSingle(img)
        img = img.resize((28, 28))
        if self.transform is not None:
            img = self.transform(img)

        return img, self.labels[idx], torch.tensor(self.group_ids[idx]), path

digitsRoot = "./digits"
handSet = HandwrittenDigits(digitsRoot, transform)

handLoader = torch.utils.data.DataLoader(handSet, batch_size=64, shuffle=False)

print("Handwritten digits found:", len(handSet))
```

```

if len(handSet) > 0:
    EvaluateWithDetails(handLoader, "Handwritten Digits", plot=True, □
    ↪printMisses=True, maxMisses=50)
else:
    print("No handwritten digits found.")

```

Handwritten digits found: 330

===== Handwritten Digits RESULTS =====

Accuracy: 0.9515

Error Rate: 0.0485

Classification Report:

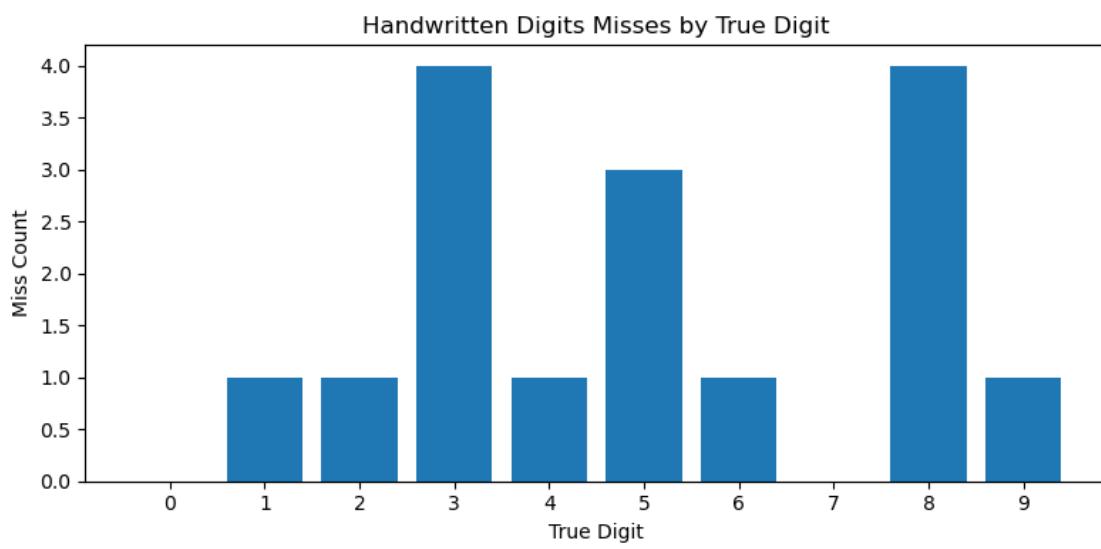
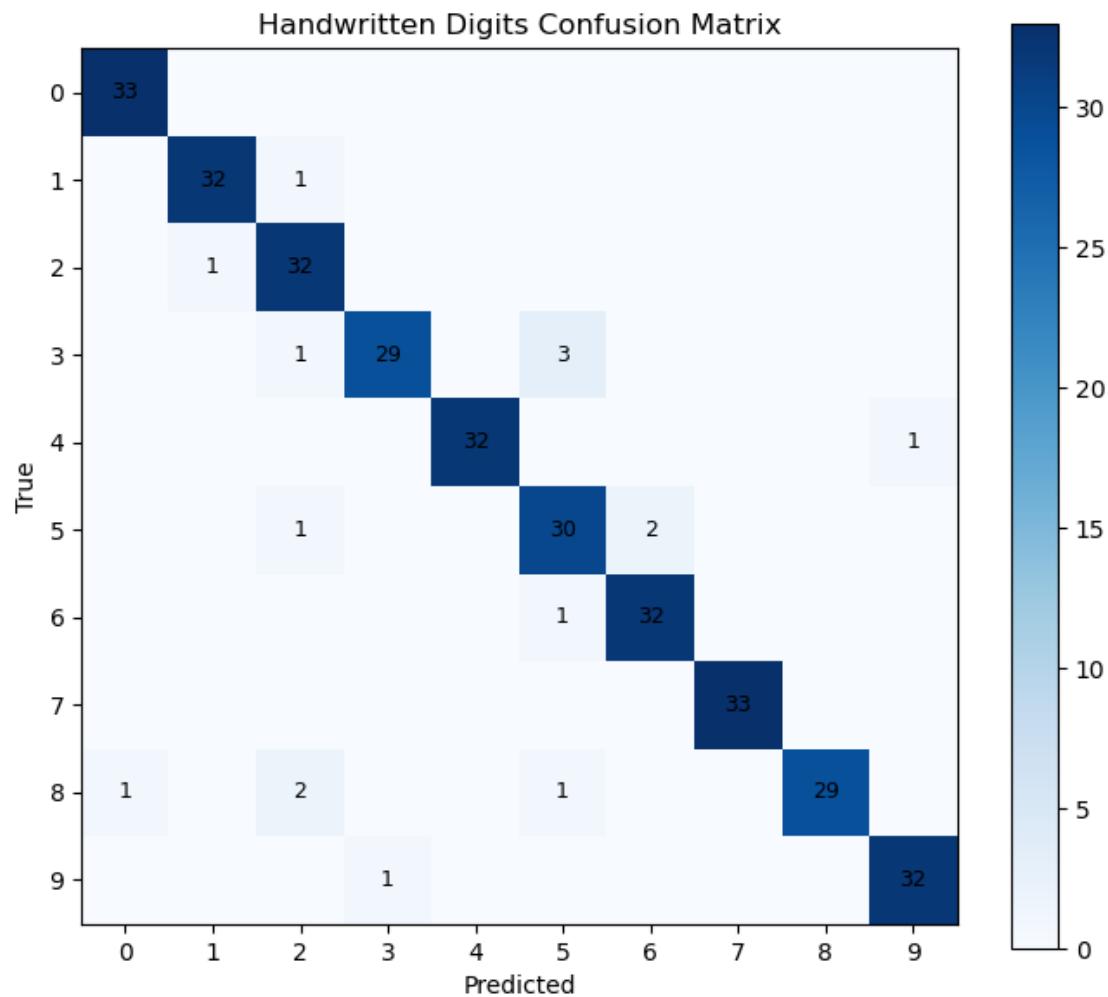
	precision	recall	f1-score	support
0	0.9706	1.0000	0.9851	33
1	0.9697	0.9697	0.9697	33
2	0.8649	0.9697	0.9143	33
3	0.9667	0.8788	0.9206	33
4	1.0000	0.9697	0.9846	33
5	0.8571	0.9091	0.8824	33
6	0.9412	0.9697	0.9552	33
7	1.0000	1.0000	1.0000	33
8	1.0000	0.8788	0.9355	33
9	0.9697	0.9697	0.9697	33
accuracy			0.9515	330
macro avg	0.9540	0.9515	0.9517	330
weighted avg	0.9540	0.9515	0.9517	330

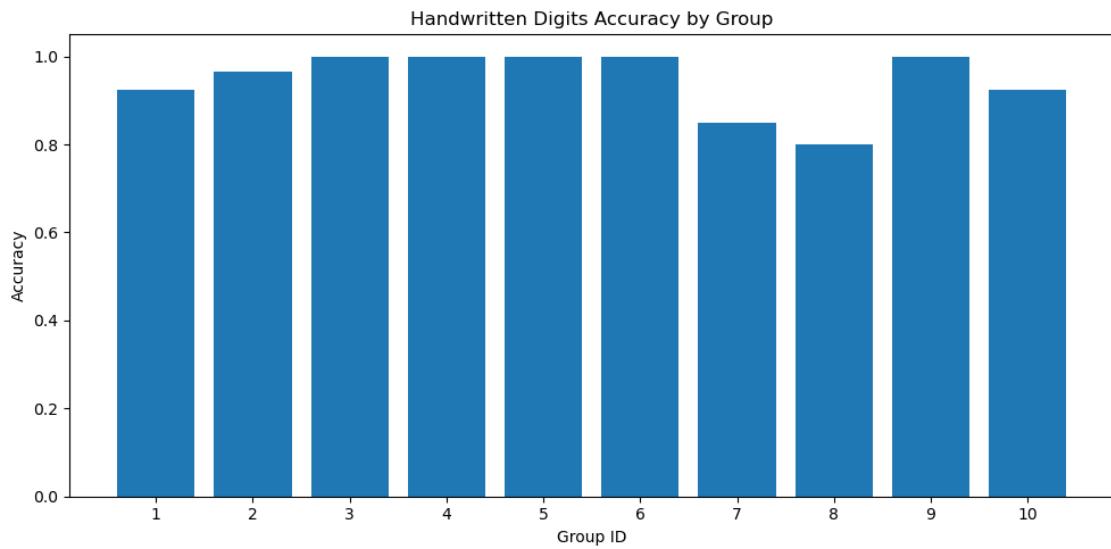
Confusion Matrix:

```

[[33  0  0  0  0  0  0  0  0  0]
 [ 0 32  1  0  0  0  0  0  0  0]
 [ 0   1 32  0  0  0  0  0  0  0]
 [ 0   0   1 29  0  3  0  0  0  0]
 [ 0   0   0 32  0  0  0  0  0  1]
 [ 0   0   1  0  0 30  2  0  0  0]
 [ 0   0   0  0  0  1 32  0  0  0]
 [ 0   0   0  0  0  0  0 33  0  0]
 [ 1   0   2  0  0  1  0  0 29  0]
 [ 0   0   0  1  0  0  0  0  0 32]]

```





Handwritten Digits mispredictions: 16 / 330

- ./digits/1-8-2.png | true=1 pred=2
- ./digits/2-1-4.png | true=2 pred=1
- ./digits/3-1-4.png | true=3 pred=5
- ./digits/3-7-1.png | true=3 pred=5
- ./digits/3-7-2.png | true=3 pred=5
- ./digits/3-8-2.png | true=3 pred=2
- ./digits/4-2-4.png | true=4 pred=9
- ./digits/5-7-2.png | true=5 pred=6
- ./digits/5-8-1.png | true=5 pred=6
- ./digits/5-8-2.png | true=5 pred=2
- ./digits/6-8-2.png | true=6 pred=5
- ./digits/8-1-1.png | true=8 pred=5
- ./digits/8-10-1.png | true=8 pred=0
- ./digits/8-10-3.png | true=8 pred=2
- ./digits/8-8-2.png | true=8 pred=2
- ./digits/9-10-3.png | true=9 pred=3